

4. Microservices & Scalability & Performance & Reliability

[Summary](#)

[Overview](#)

[Key guidelines](#)

[Example separation into microservices](#)

Summary [🔗](#)

- Decompose the system into independent, cloud-native services (e.g. auth, payments, communication, bookkeeping) and use gRPC/HTTP plus message queues for inter-service calls.
- Apply resiliency patterns (retries, circuit breakers, fallbacks), define SLAs/SLOs/SLIs, and ensure redundancy to guarantee reliability.
- Leverage platform autoscaling (KEDA), load balancing, API gateway/service discovery, and caching layers for performance and scalability.

Pending Decisions

- Select and configure the API gateway/service discovery solution (Azure API Management vs Application Gateway vs custom).
- Finalize the data consistency pattern (event sourcing, sagas, single-service DB) and caching strategy.

Open Items

- Document full microservice boundaries and integration workflows beyond the initial examples.
- Establish CI/CD and deployment patterns (e.g. blue-green/ Canary releases).
- Clarify security boundaries, token management, and RBAC between services.
- Review and align with detailed integration requirements with the other topics.

Overview [🔗](#)

A microservices-architected distributed system ensures scalability, reliability, and performance by breaking down the system into smaller, independent services. Each service should follow industry best practices, prioritize resiliency, and leverage cloud-native solutions.

Key guidelines [🔗](#)

Trust the lead developer / tech lead and overall team experience for application architecture - with secondary knowledge boost brought in through AI

- **Place a high importance on inter-service communication:**
 - consider gRPC for efficient communication between services, especially for high-performance needs. Use HTTP for external APIs where compatibility is more important
 - implement message queues for asynchronous communication and background processes (even one file's download should be queued up in most cases, especially if it takes away resources from a front-end application)
 - always keep in mind data consistency questions and avoid designing solutions on platforms which are not a good fit for such purpose (*cough cough* Kafka with equal-hierarchy topics)

- **Resiliency & reliability:**

- design services to handle failures gracefully using retry policies, circuit breakers, and fallback mechanisms.
 - for example [Dapr - Distributed Application Runtime](#) ([Microservice APIs powered by Dapr | Microsoft Learn](#)) and inbuilt runtime specific APIs ([Build resilient HTTP apps: Key development patterns - .NET | Microsoft Learn](#))
- redundancy: ensure critical resources have a degree of redundancy to minimize downtime in unforeseen events
- try to provide SLAs, SLOs and SLIs (promise, objective, measurement): to measure and guarantee reliability
- consider short-circuit strategies to isolate failing services, preventing cascading failures

- **Scalability:**

- prefer platform-provided auto scaling options
 - like [KEDA](#) (used natively also by various cloud platforms, e.g. [Scaling in Azure Container Apps | Microsoft Learn](#))
- load balancing, gateway and routing options: again, *prefer platform-provided options over coupling such things into core business APIs*:
 - which can also handle rate limiting ([Rate Limiting pattern - Azure Architecture Center | Microsoft Learn](#)), throttling, caching, traffic distributing, etc..
 - load balancing happens on various layers, keep in mind. See [Load-balancing options - Azure Architecture Center | Microsoft Learn](#)
 - [Azure API Management - Overview and key concepts | Microsoft Learn](#) & [What is Azure Application Gateway | Microsoft Learn](#)
- authentication & authorization: again, split it out of the core as much as possible
 - see also [1. Authentication & Authorization](#)

- **Data management:**

- when using SQL databases, ensure only one application owns and manages it to maintain consistency and enable proper scalability
- utilize various caching layers, with cloud-native approach (emphasis on distributed systems)

- **Cloud native approach:**

- see [16. Hosting](#), especially understand implicit dependencies and reliance on third-party and that the cloud is a shared responsibility
- think about service discovery, cataloging (e.g. [Azure API Center - Overview - Azure API Center | Microsoft Learn](#))

Example separation into microservices [🔗](#)

- Payment gateway microservice → docs started at: [LISASPORTS/new-payment-gateway](#)
- Communication microservice → [LISASPORTS/communication-service: A centralized communication service for Lisa](#)
- Auth service → [LISASPORTS/OAuthServer](#), see also [1. Authentication & Authorization | OAuthServer Responsibilities](#)
- Bookkeeping system microservice
- see also further integrations requirements at [6. Communication & Notifications & Integrations](#)